TPC5

December 3, 2018

In [23]:
```
import numpy as np


print("1a)")

Qt = 3.08
c = 1.0

#(4, -b r)
Q4r = np.array([
    [3.23 , 3.38, 3.25, 3.22]
])

#(5, -b r)
Q5r = np.array([
    [3.08, 3.25, 3.57, 3.22]
])

Qnew = Qt + 0.1 * (c + 0.99 * np.min(Q5r) - Qt)


print("Q-Learning Q-value = %s" % Qnew)

print("\n1b)")

Qt = 3.08

Qnew = Qt + 0.1 * (c + 0.99*Q5r[0][3] - Qt)

print("Sarsa Q-value = %s" % Qnew)


print("\n1c)\n")

print("Off-policy refers to reinforcement learning methods")
print("that learn the value of a policy while following another, like Q-learning.")
print("On-policy refers to reinforcement learning methods that learn the value of ")
print("the policy that the agent is following.")
```

```
        print("SARSA is more stable than Q-learning, as long as ")
        print("learning policy changes smoothly with Qt")
        print("The difference between the two update computations: ")
        print("the Q-learning chooses the minimum value of the Q-values ")
        print("for each action, in this code np.min(Q5r).")
        print("While the SARSA algorithm it is based on the the policy it")
        print(" is following, in this case coded Q5r[0][3].")

1a)
Q-Learning Q-value = 3.17692

1b)
Sarsa Q-value = 3.19078

1c)

Off-policy refers to reinforcement learning methods
that learn the value of a policy while following another, like Q-learning.
On-policy refers to reinforcement learning methods that learn the value of
the policy that the agent is following.
SARSA is more stable than Q-learning, as long as
learning policy changes smoothly with Qt
The difference between the two update computations:
the Q-learning chooses the minimum value of the Q-values
for each action, in this code np.min(Q5r).
While the SARSA algorithm it is based on the the policy it
 is following, in this case coded Q5r[0][3].
```